



## Fourth Industrial Summer School

Advanced Machine Learning

Neural Networks and Deep learning-Part4

# Session Objectives

- ✓ Introduction
- ✓ Fundamentals
- ✓ Neural Network Intuitions
- ✓ 2-Layer Neural Network
- ✓ Deep Neural Networks
- ✓ CNNs
- ✓ Keras with Tensorflow



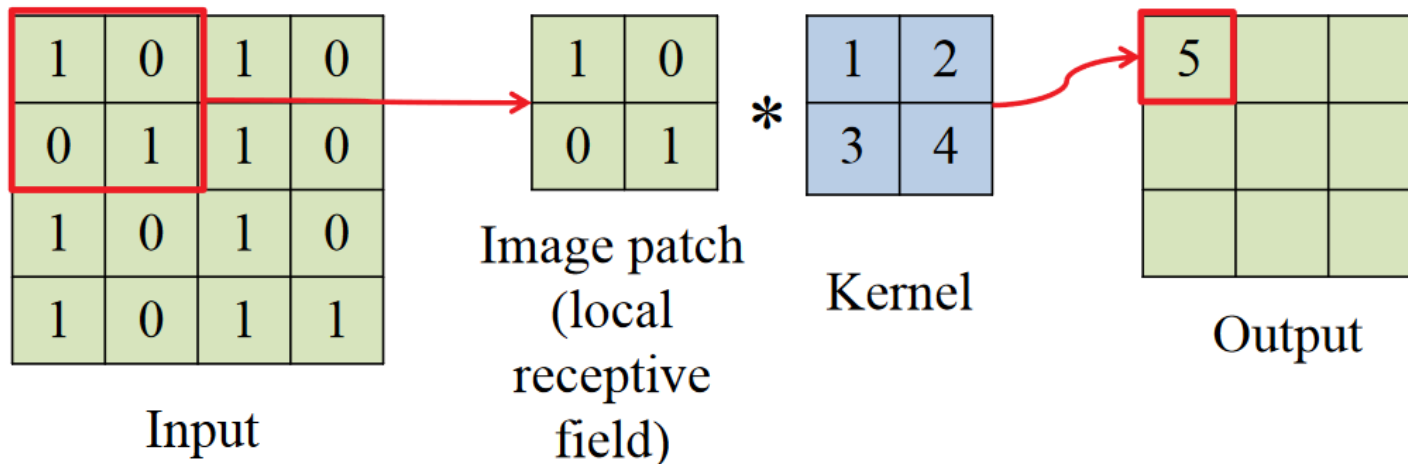
# CNNs

# Why CNNs

- Used mostly with images
- Translation invariance
- How was it before NNs?

CNN:

- A dot product of a kernel (aka filter) and a patch of an image (local receptive field) of the same size.



# CNNs-Edge detection



Original  
image

Kernel

-1	-1	-1
-1	8	-1
-1	-1	-1

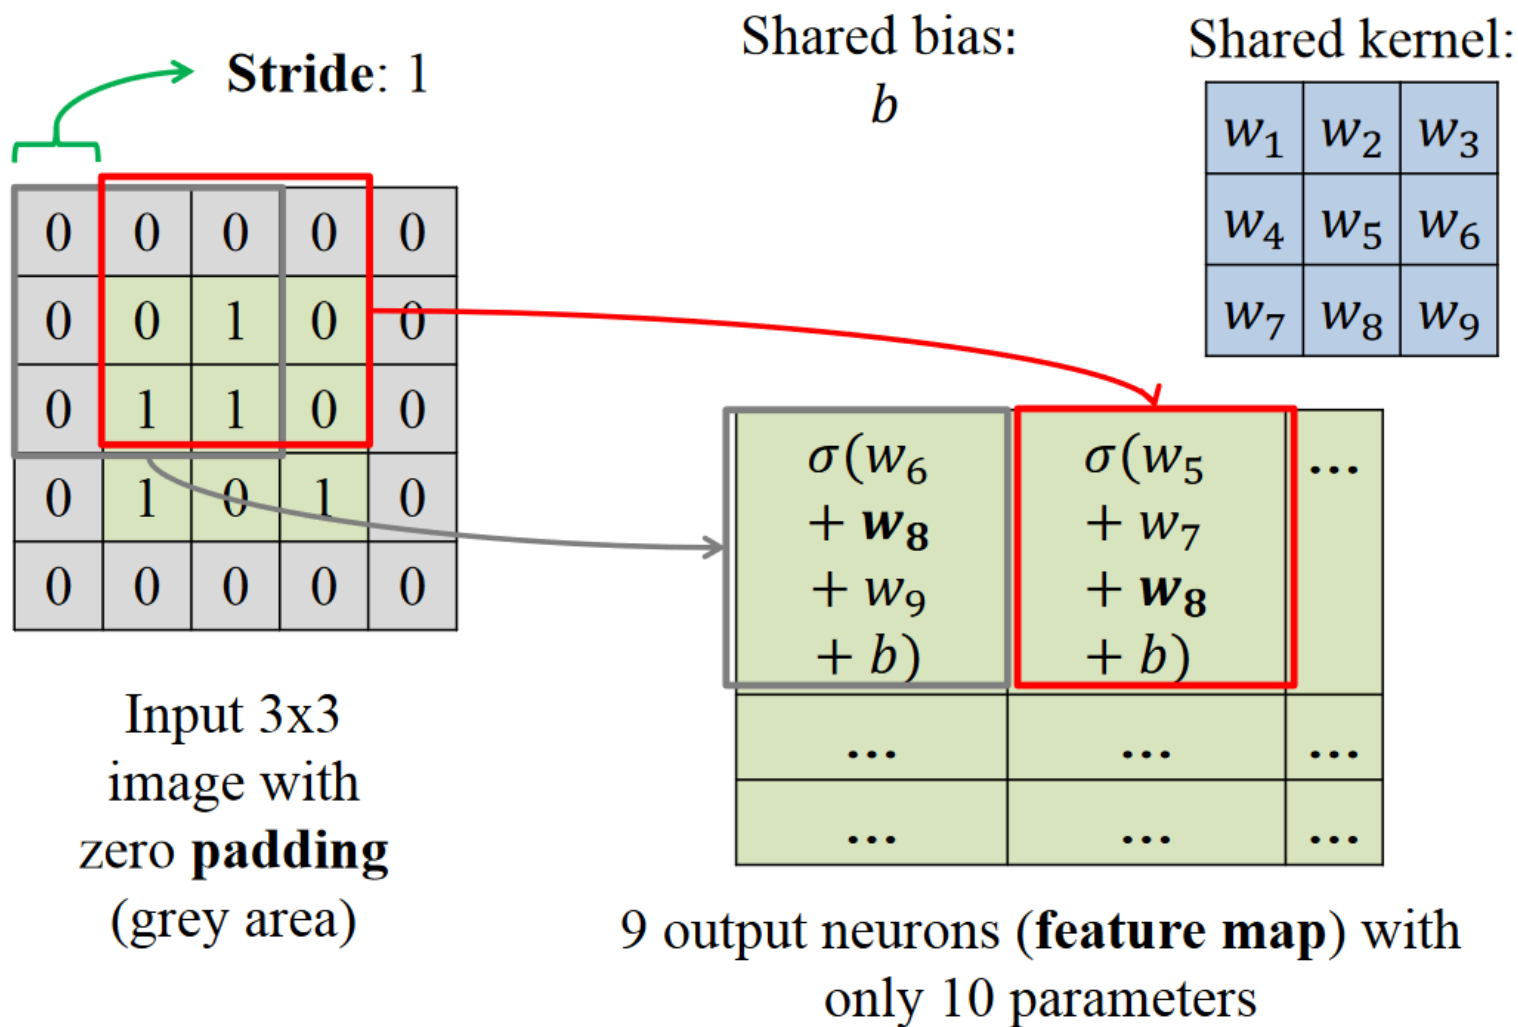
\*      =      Edge detection

A small, square image showing the edge detection result of the squirrel image. The image is mostly black, with white lines highlighting the edges of the squirrel's head and features.

Sums up to 0 (black color)  
when the patch is a solid fill

- It can also do sharpening and blurring.

# CNN-At work



# Backpropagation for CNN



- Sum of all the gradients used by a weight in a layer.


# Other properties



- Same kernel is used for every output neuron of a layer, thus fewer weights to be trained.
- So, 30x30 image with 5x5 kernel will have how many parameters?
- What about fully connected layer?
- Very important in high-dimension images.




# CNN for colored images



# Stacking CNNs


- To capture region of interest of different resolutions
- One kernel is not enough.



# What about translation of regions

- Pooling solves it a greater degree.
- Pooling layer does not have a kernel
  - It calculates maximum (or average) over the patch.


- Backpropagation?



2x2 **max pooling** with stride 2

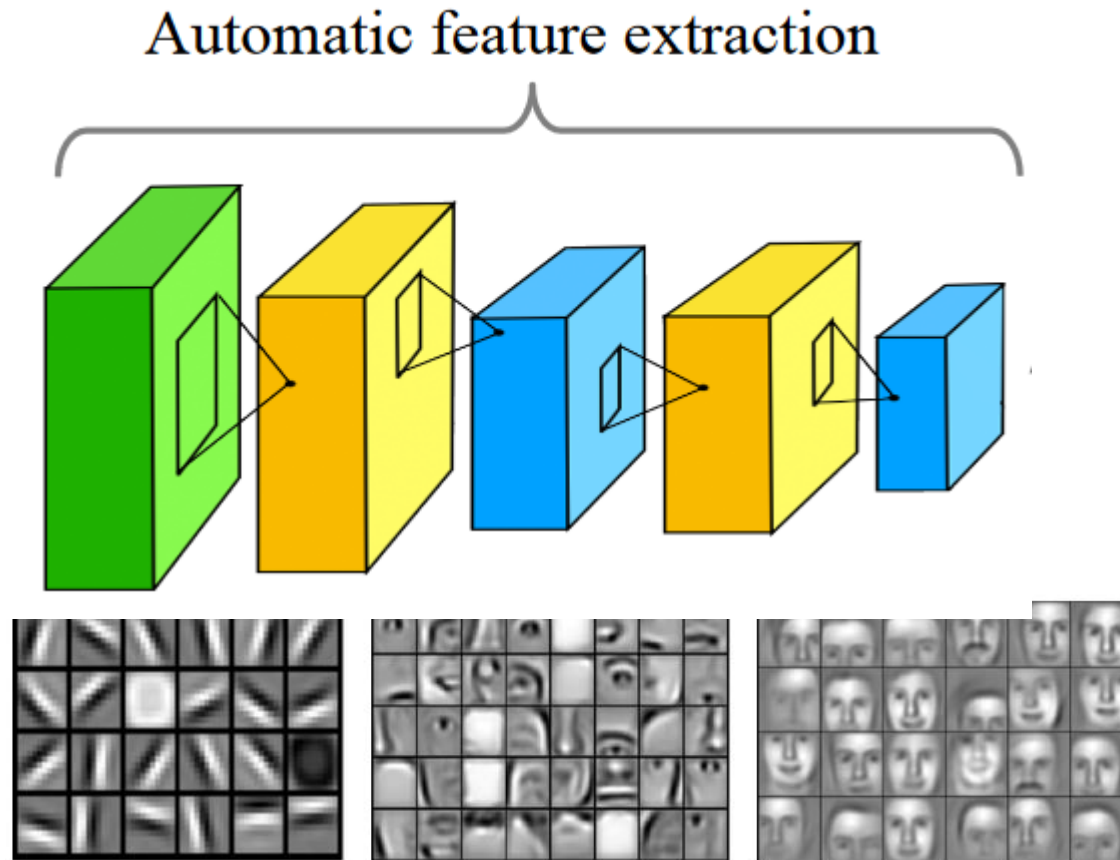
# An example CNN

## ■ LeNet-5 architecture



<http://yann.lecun.com/exdb/publis/pdf/lecun-98.pdf>

# Learning deep representations



<http://web.eecs.umich.edu/~honglak/icml09-ConvolutionalDeepBeliefNetworks.pdf>

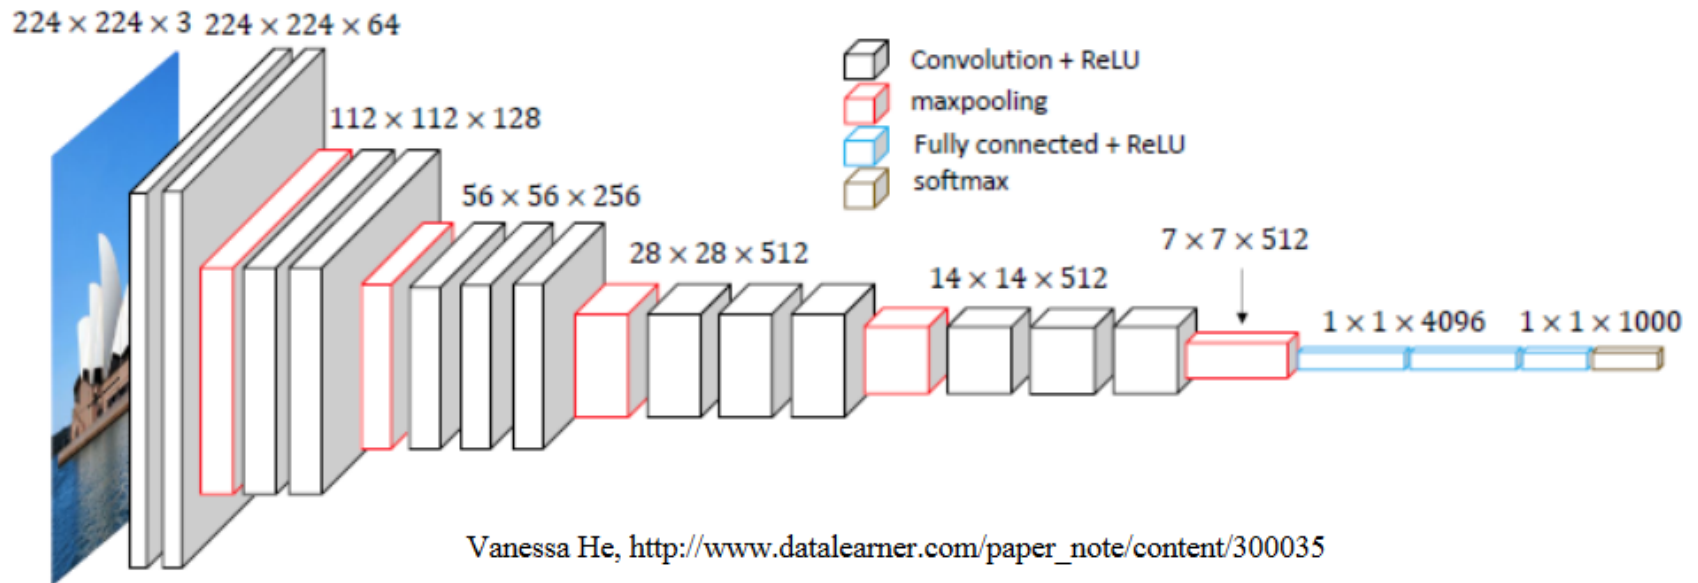
# Optimization, Hyper-parameters, etc.



- Learning rate
- Activation function
- Weight initialization
- Batch normalization
- Dropout
- Data augmentation

4IR Summer School


# VGG



- 138 million parameters
- Training: 2-3 week (4 GPUs)



# ResNet



Kaiming He, <https://arxiv.org/pdf/1512.03385.pdf>

- 152 layers!!!
- 7x7 convolutional layers, 3x3 convolutional layers, batch normalization, max and average pooling.
- Parameters: 60 million
- Training time: 2-3 weeks (8 GPUs)


# Transfer Learning



- CNNs as feature extractor
- Use number of layers based on the similarity of the data/task
- Other layers can be used for fine-tuning
- Different scenarios based on data size and similarity

# Other tasks involving CNNs

## ■ Image segmentation



Long, Shelhamer, and Darrell, "Fully Convolutional Networks for Semantic Segmentation", CVPR 2015

## ■ Object localization

# Keras

with TensorFlow

# Keras



A high-level library which can work over tensorflow

- Quick model development
- User friendliness
- Modularity
- Easy extensibility
- Work with Python

# Keras Models



- Sequential: A linear stack of layers
- Model (Functional API): Any arbitrary setup

# Keras Sequential

- A linear stack of layers

```
from keras.models import Sequential
from keras.layers import Dense, Activation

model = Sequential([
    Dense(32, input_shape=(784,)),
    Activation('relu'),
    Dense(10),
    Activation('softmax'),
])
```

```
model = Sequential()
model.add(Dense(32, input_shape=(784,)))
```

Note!!! With tensorflow implementation of keras, we need to use tensorflow.keras instead of keras:

e.g., from tensorflow.keras.model import Sequential

# model.add()

- Stacking layers in sequence

---

```
from keras.layers import Dense

model.add(Dense(units=64, activation='relu', input_dim=100))
model.add(Dense(units=10, activation='softmax'))
```

## Core Layers:

- Dense: Fully connected layers
- Activation
- Dropout
- Flatten



# Layers- Conv2D

- 2D convolution layer (e.g. spatial convolution over images).

## Important Arguments

- Filters: No of filters (the dimensionality of the output space)
- `kernel_size`
- Strides
- Activation
- `kernel_initializer`

# Layers-MaxPooling2D



- Max pooling operation for spatial data.

Important Arguments:

- `pool_size`
- `strides`

# model.compile()

- Once the model structure is ready, configure its learning process with:

```
# For a multi-class classification problem  
model.compile(optimizer='rmsprop',  
              loss='categorical_crossentropy',  
              metrics=['accuracy'])
```

- <https://keras.io/optimizers/>
- <https://keras.io/losses/>

```
# For a binary classification problem  
model.compile(optimizer='rmsprop',  
              loss='binary_crossentropy',  
              metrics=['accuracy'])
```

```
# For a mean squared error regression problem  
model.compile(optimizer='rmsprop',  
              loss='mse')
```

```
model.compile(loss=keras.losses.categorical_crossentropy,  
              optimizer=keras.optimizers.SGD(lr=0.01, momentum=0.9, nesterov=True))
```

# model.fit()

Perform the actual training, it has following main arguments:

- Xs and Ys
- batch\_size
- Epochs
- validation\_split

```
# x_train and y_train are Numpy arrays --just like in the Scikit-Learn API.  
model.fit(x_train, y_train, epochs=5, batch_size=32)
```

# model.evaluate()



- Evaluate the performance

```
loss_and_metrics = model.evaluate(x_test, y_test, batch_size=128)
```

# Keras Models-Functional API

- The Keras functional API is the way to go for defining complex models:
  - Multi-output models
  - Models with shared layers

```
from keras.layers import Input, Dense
from keras.models import Model

# This returns a tensor
inputs = Input(shape=(784,))

# a Layer instance is callable on a tensor, and returns a tensor
x = Dense(64, activation='relu')(inputs)
x = Dense(64, activation='relu')(x)
predictions = Dense(10, activation='softmax')(x)

# This creates a model that includes
# the Input Layer and three Dense Layers
model = Model(inputs=inputs, outputs=predictions)
model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
model.fit(data, labels) # starts training
```

# References

- <https://keras.io>
- Introduction to Deep Learning, National Research University Higher School of Economics
- Fei-Fei Li Convolutional Neural Networks for Visual Recognition, Stanford University (<http://cs231n.stanford.edu/>)
- Michael A. Nielsen, "Neural Networks and Deep Learning", Determination Press, 2015
- Andrew Ng, Neural Networks and Deep Learning, Stanford University
- <http://web.eecs.umich.edu/~honglak/icml09-ConvolutionalDeepBeliefNetworks.pdf>
- Andrew Ng, Machine Learning Yearning, [deeplearning.ai](http://deeplearning.ai)